Table of Contents

An Accelerator to Provision Databricks with Azure DevOps

Objective

The purpose of this tutorial is a walkthrough to how one could securely set up Azure Databricks with an Azure Storage container mounted. Azure DevOps allows for an easy integration between all the different services required. Below is a visual representation of the overall architecture of the tutorial.

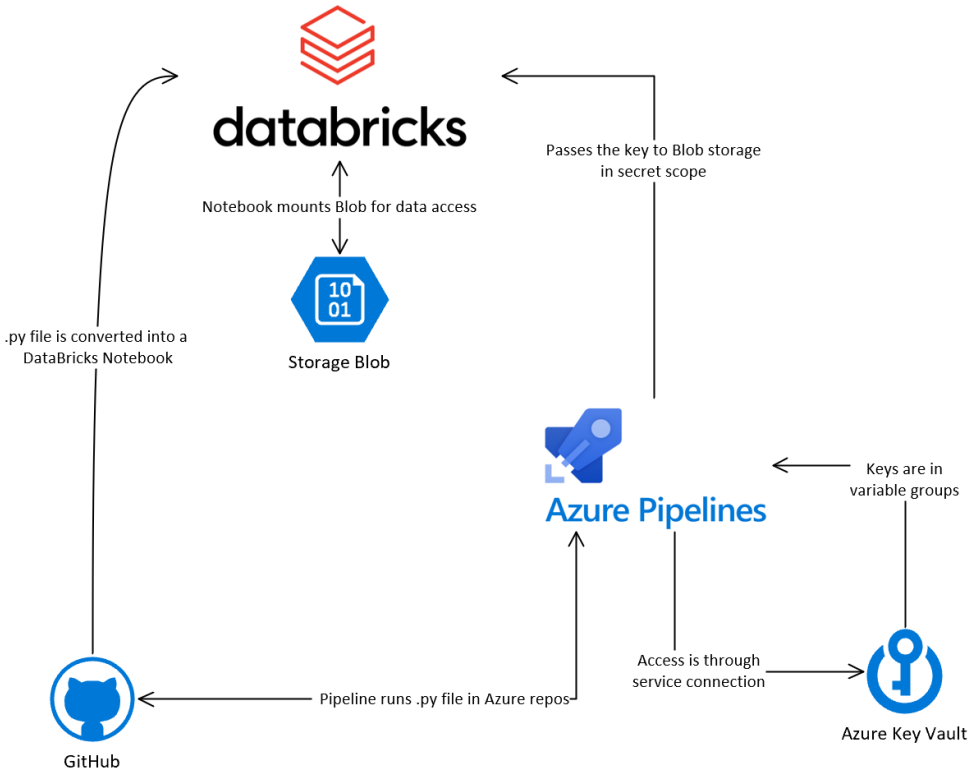

Prerequisites

For this project, the following prerequisites are assumed:

- Azure Active Directory Tenant (As in, can access Azure Portal)
- Azure Subscription
- GitHub account

Getting Started

For this section, the following resources should be created:

- Resource Group
- Storage Account
- Container in the storage account
- Key Vault with a secret
- Azure Databricks Service

Note: It is highly advisable to use the same names, but the user could choose different names. However, if the names are changed, make sure to change the respective names in right sections.

Resource	Name	Should be a different name?
Resource Group	databricks-provisioning-accelerator	NO
Storage Account	storageaccelerator	YES
Container in storage	storage-container	NO
Key Vault	acceleratorkeyvault	YES
Azure Databricks Service	accelerator-databricks	NO

Creating a Resource Group

1. Login into http://portal.azure.com/
2. In the search bar, search for **Resource Groups**
3. Click the *+ ADD* button
4. Fill in the required fields (as shown below)
5. Click *Review + Create*
6. Click *Create*

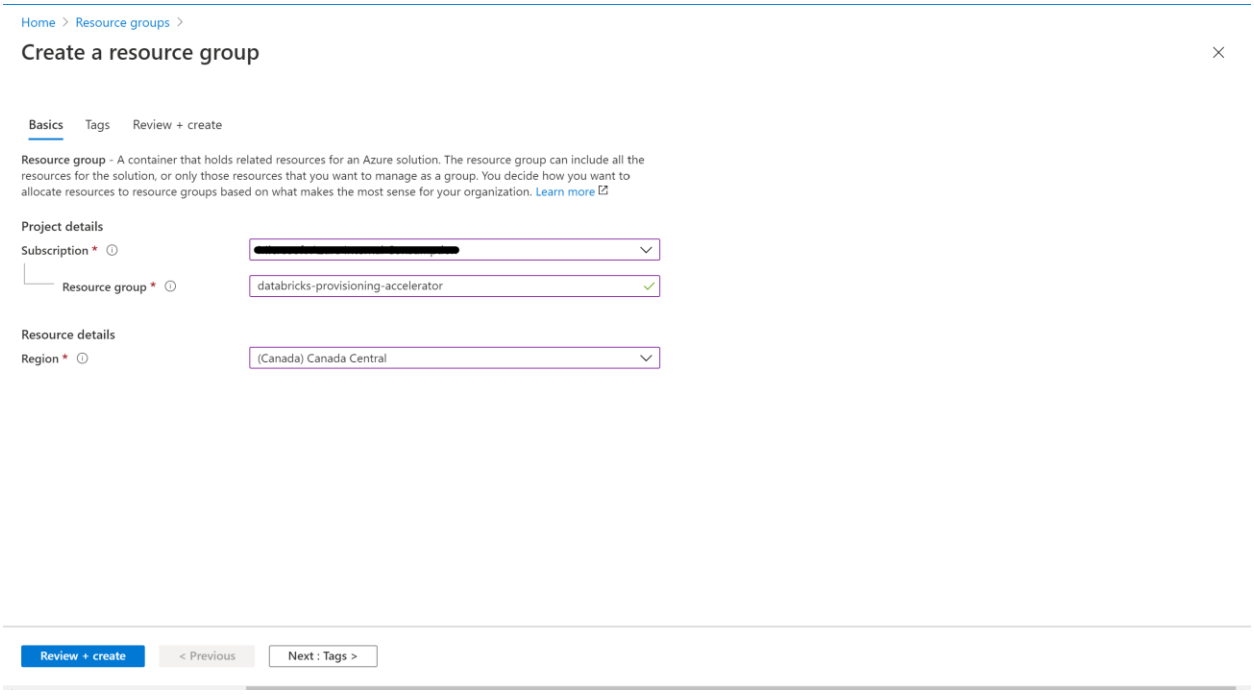

Creating a Storage Account

(Note: One could also provision Azure Data Lake Gen 2 (with hierarchical namespace enabled))

1. Open the Resource Group created previously
2. Click *+ ADD* (as shown below)

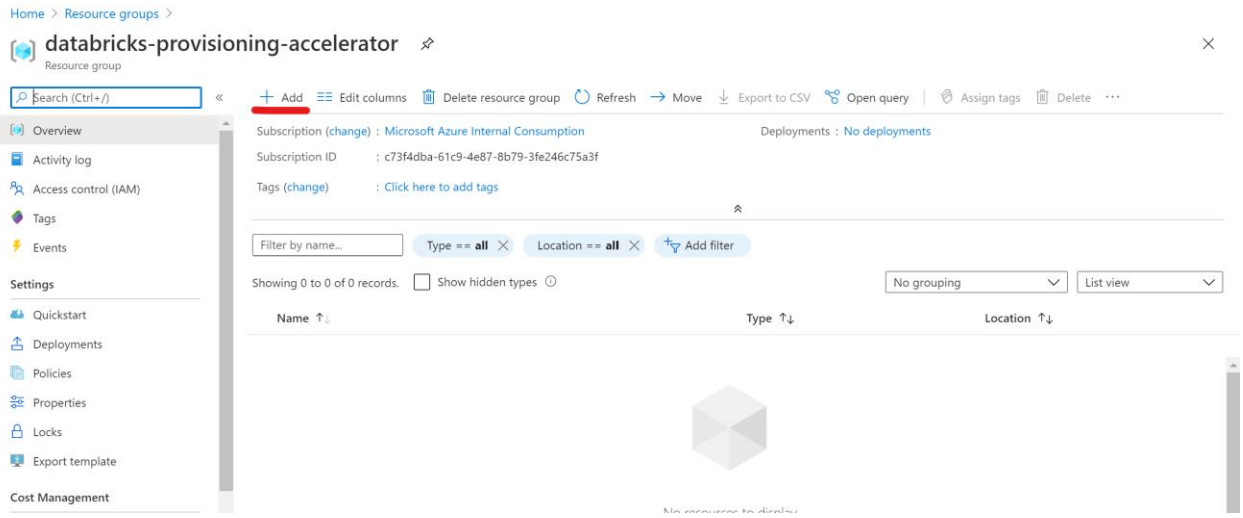

3. Search for **Storage Account** and click *Create*
4. Fill in the fields as required and click *Review + Create*
5. Click *Create*

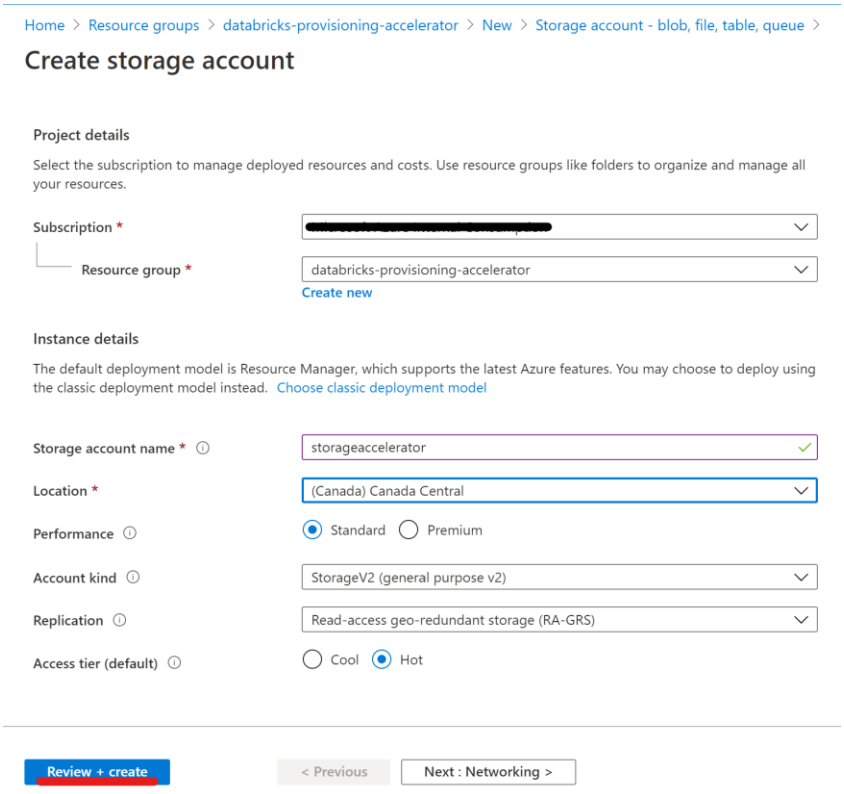

Creating a Container in a Storage Account

1. Open the Resource Group created earlier
2. Click on the storage account created in the previous step
3. Click on *Containers*

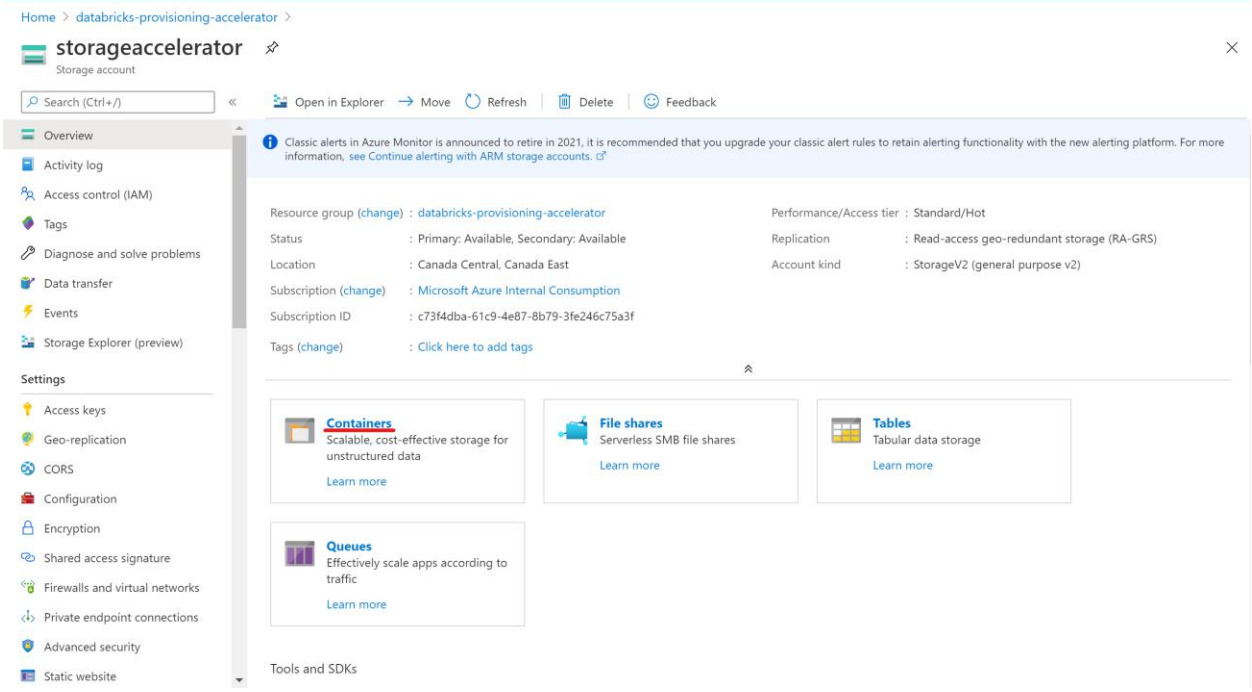

4. Click on *+ Container*
5. Fill the *Name* field and click *Create*

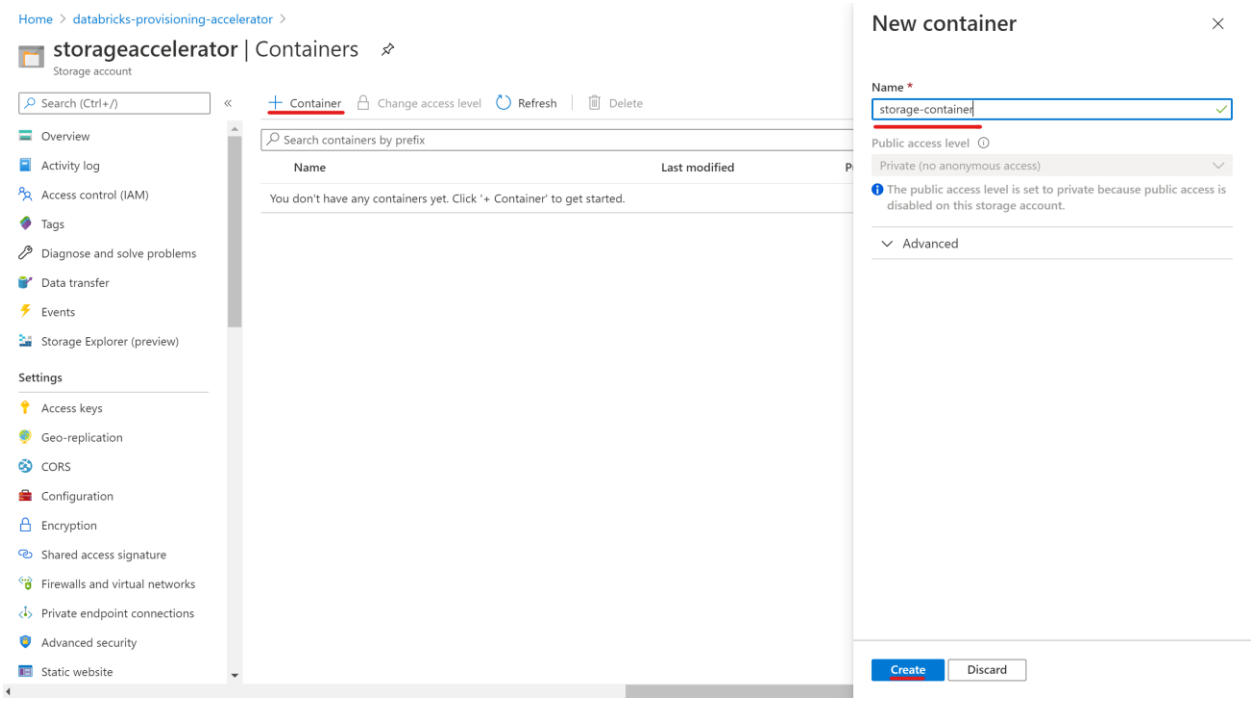

Creating a Key Vault account with a Secret

For more info click here

1. Open the Resource Group created earlier
2. Click *+ ADD*
3. Search for **Key Vault** and click *Create*
4. Fill in the fields as required and click *Review + Create*

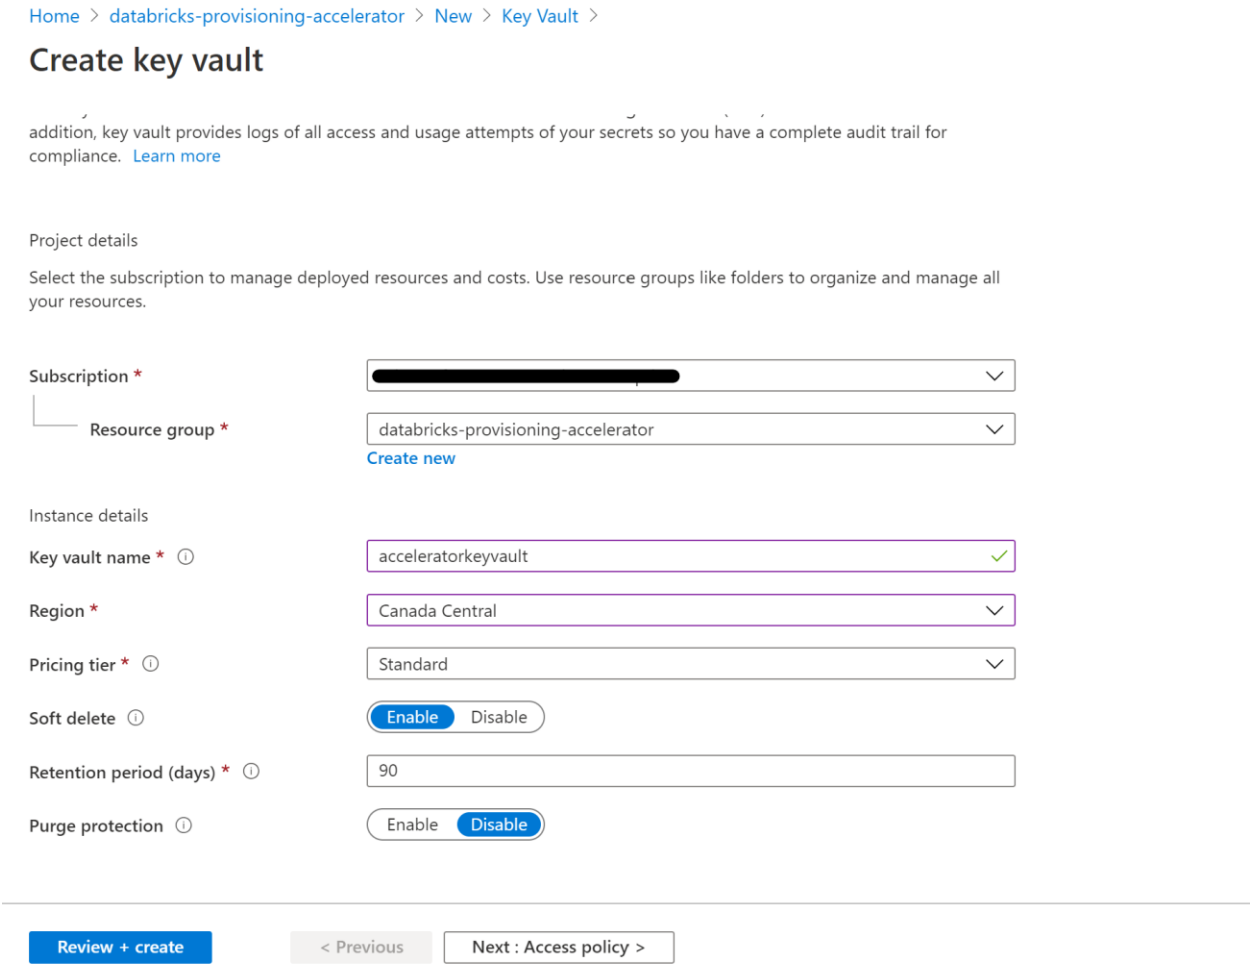

5. Click *Create*
6. Go to the Resource Group and open *acceleratorkeyvault*
7. On the left menu, under Settings, click *Secrets*
8. Click *+ Generate/Import*
9. Fill in the fields.
    a. **Note:** the key must be **accelerator-storage-key**, otherwise it will need to be changed in the pipeline's YAML file
    b. *Value* is retrieved from *storageaccelerator* > *Access keys* > copy *key1*
    c. For more info click here
10. Click *Create*

Creating an Azure Databricks Service

1. Open the Resource Group created earlier
2. Click *+ ADD*
3. Search for **Azure Databricks** and click *Create*
4. Fill in the fields. Note that the *Pricing Tier* should be *Premium*

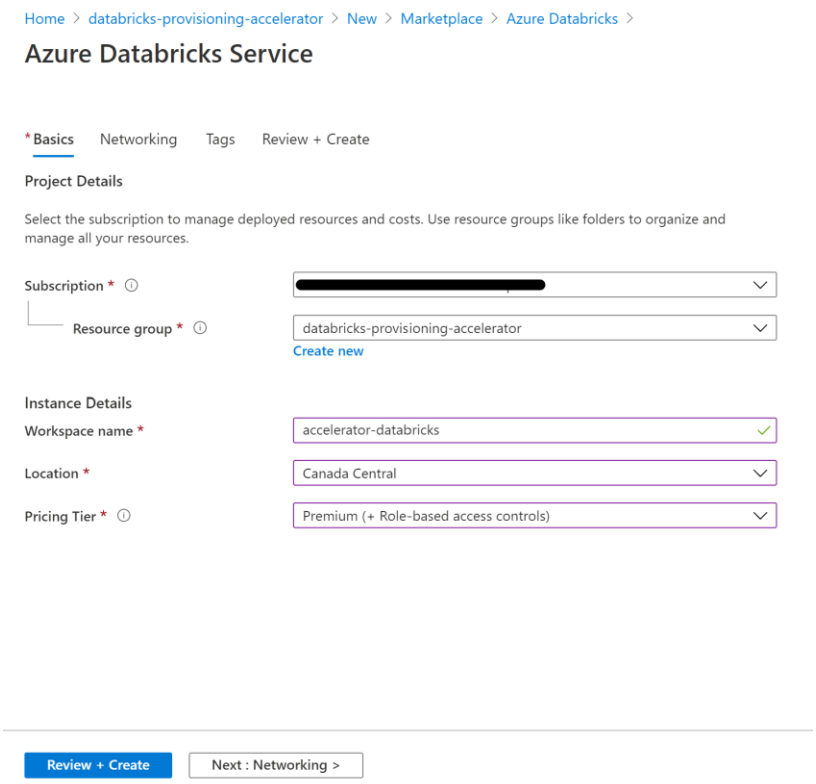

5. Click on *Review + Create*
6. Click *Create*

Steps to Set-up Databricks

1. Creating an App registration

For more information, click here.

1. Sign in to the Azure portal
2. Search for and select *Azure Active Directory*
3. Select *App Registrations*
4. Select *New Registration*
5. Fill the required fields as shown below and click *Register*

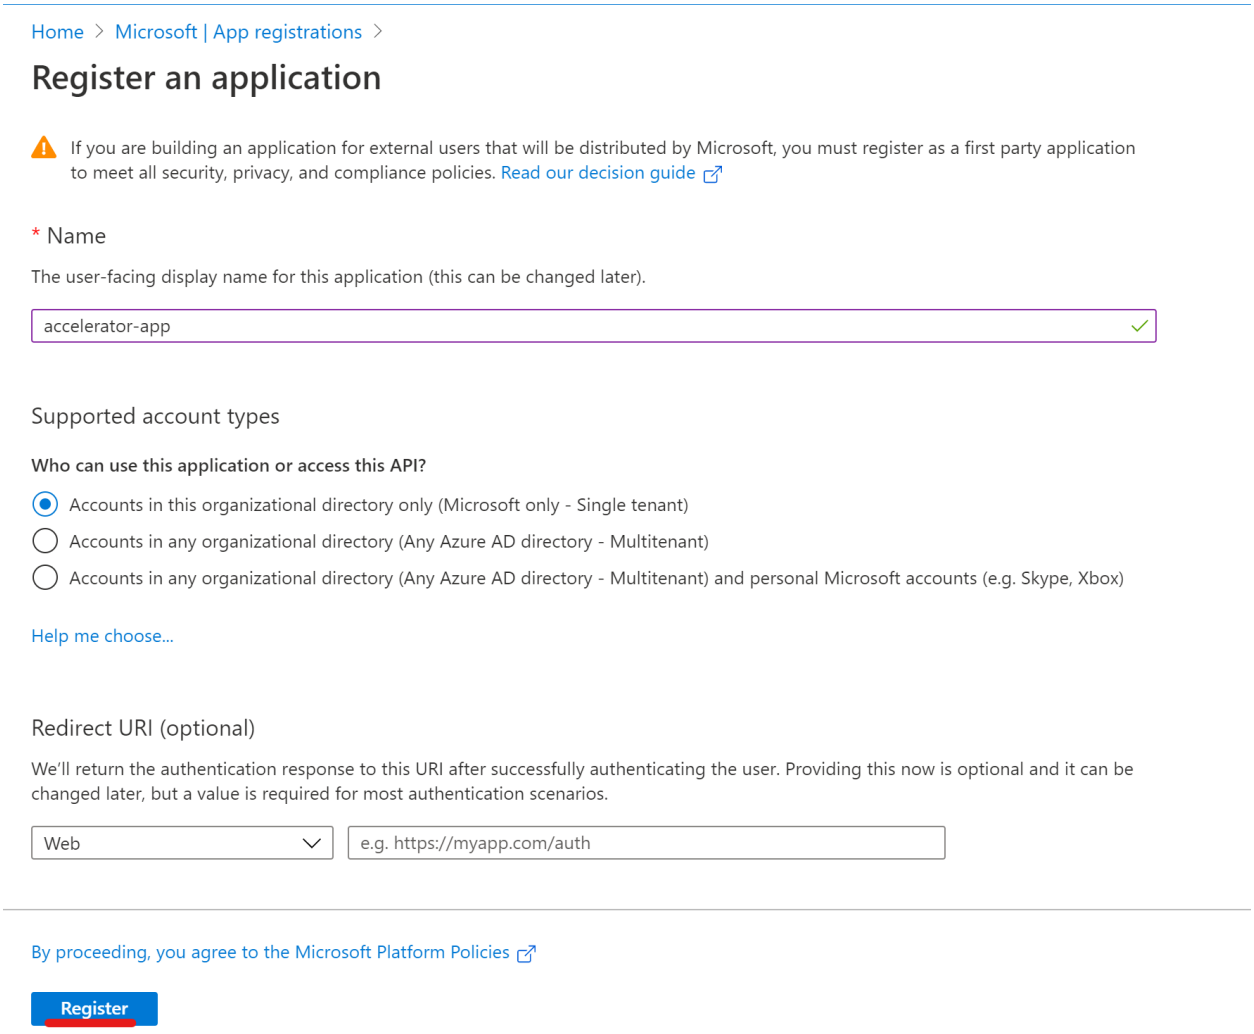

2. Assign a role to the application

For more information, click here.

1. In Azure portal, search for and select *Subscriptions*
2. Select the subscription you wish to use
3. Select *Access control (IAM)*
4. Select *Add Role Assignment*
5. Set the Role as *Contributor.* (Note: this is a high-privileged setting, one can give the service principal access to the resource group or even to the individual services.)
6. Select the application created previously and click *Save*

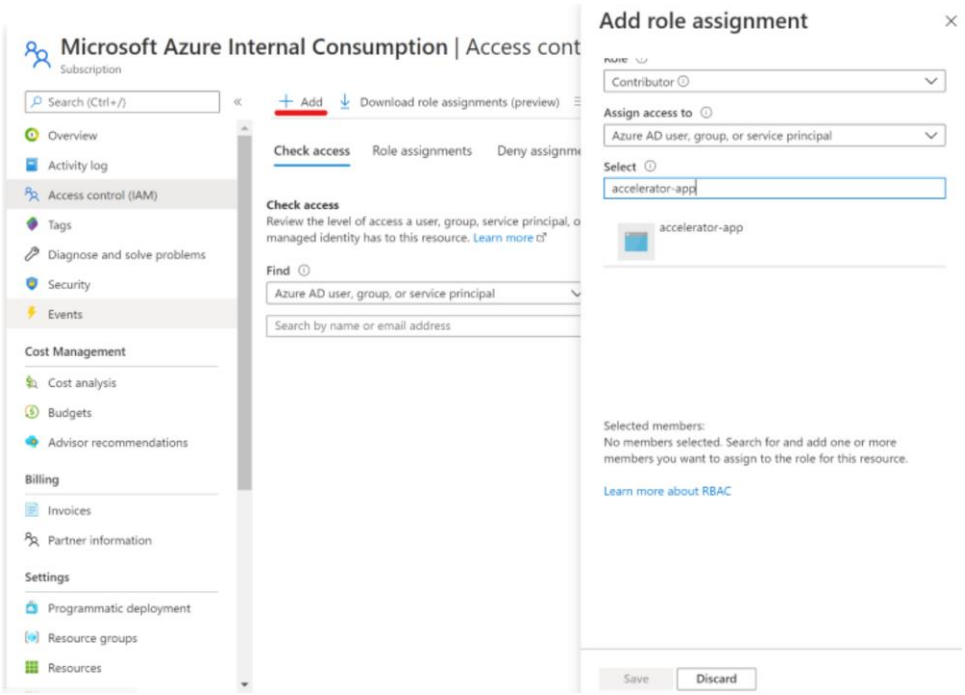

3. Get the tenant and app ID values

We need to retrieve the tenant and app ID values for the next step. Therefore, store them somewhere safe, as they are needed later. For more information, click here.

1. In Azure portal, search and select *App Registrations*
2. Select the app created earlier
3. Copy the *Application (client) ID* and then the *Directory (tenant) ID*

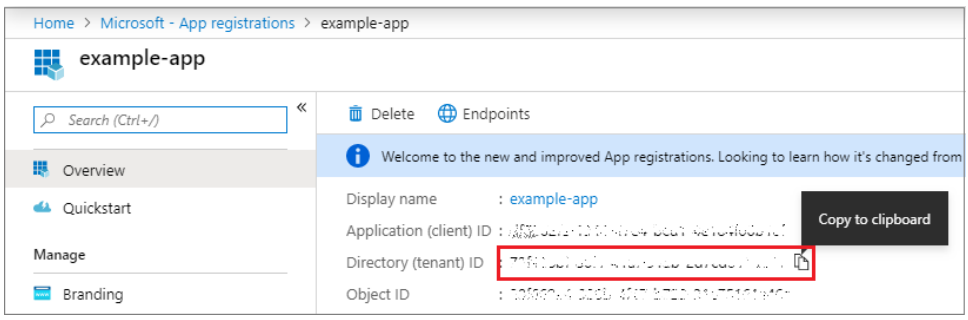

4. Creating an application secret

1. In Azure portal, search and select *App Registrations*
2. Select the app created earlier
3. Navigate to *Certificates & secrets*
4. Click *+ New client secret*
5. Fill in the required fields
6. Click *Add*
7. **Copy the value of the secret and store it securely (it can only be seen once)**

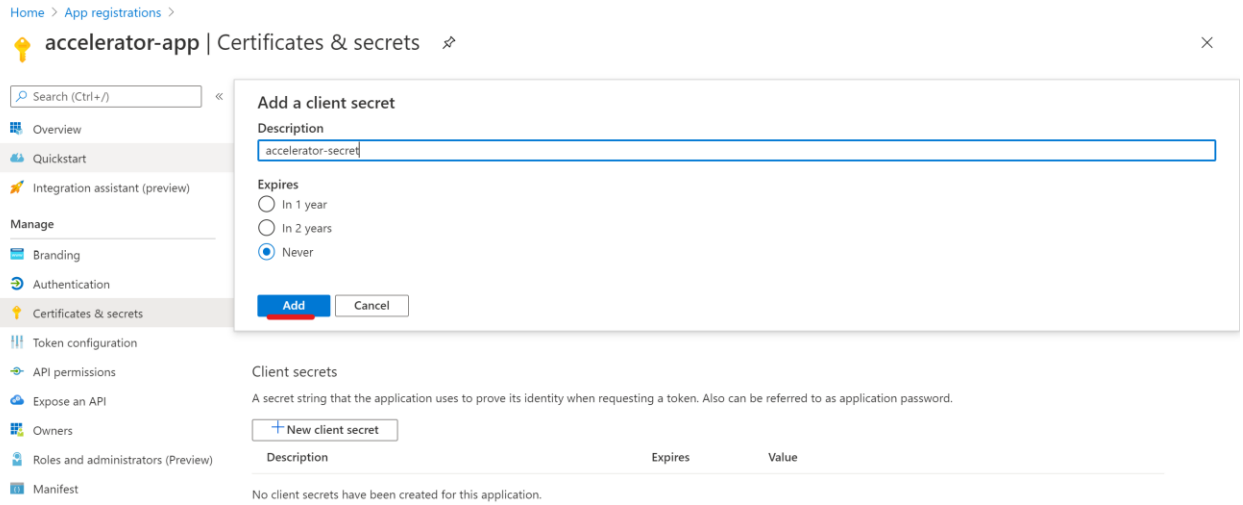

5. Retrieve Subscription Name and ID

1. In Azure portal, search and select *Subscriptions*
2. Select the subscription used throughout this project
3. Copy the subscription's name and subscription ID and store them safely

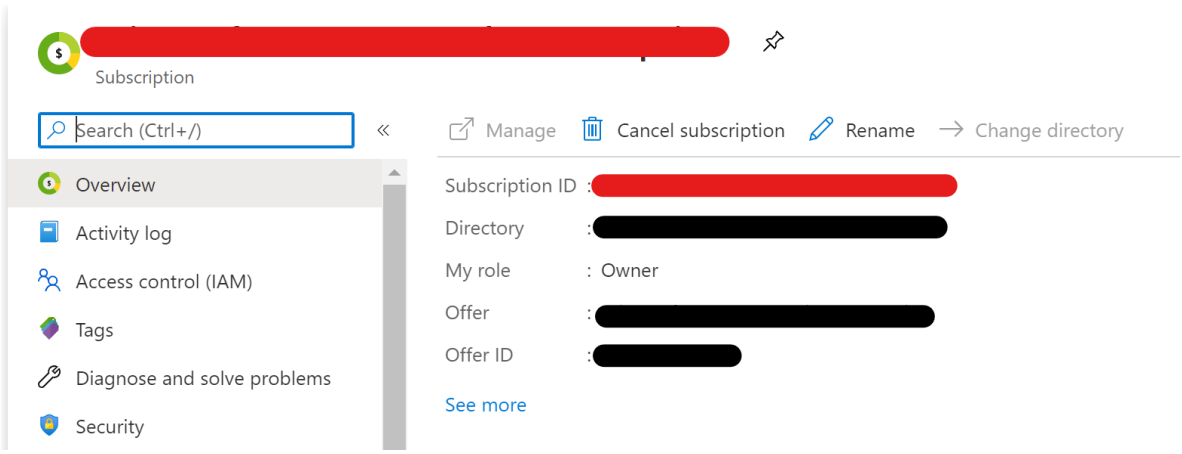

By now you should have the following hash values (necessary for the next step):

- Application (client) ID
- Directory (tenant) ID
- Value of the app's secret
- Subscription name and ID

6. Creating an Azure DevOps Project and a Service Connection

1. Login into https://dev.azure.com/
2. Click *+ Create Project*
3. Fill in the required fields and click *Create*
4. At the bottom of left menu bar, click on the gear logo for the Project Settings
5. Under *Pipelines*, click *Service Connections*
6. Click on *New service connection*
7. Click on *Azure Resource Manager*
8. Click on *Service principal (manual)*
9. Paste the Subscription ID and Name under *Subscription Id* and *Subscription Name* fields
10. Paste the Application (client) ID under the *Service Principal Id* field
11. Paste the value of the app's secret under *Service principal key*
12. Paste the Directory (tenant) ID under the *Tenant ID*
13. Click *Verify*, it's necessary to have a Verification Succeeded response
14. Fill in the *Service connection name* field
15. Click *Verify and save*

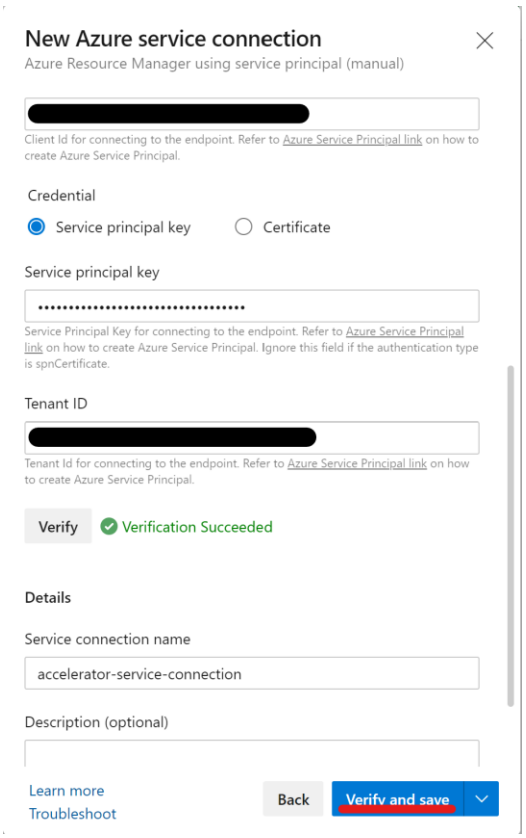

7. Forking the GitHub Repository

1. Navigate and sign in to https://github.com/
2. Navigate to https://github.com/waddahAldrobi/azure-devops-databricks-accelerator
3. Click on *Fork* to fork the repository on the master branch (as shown below)

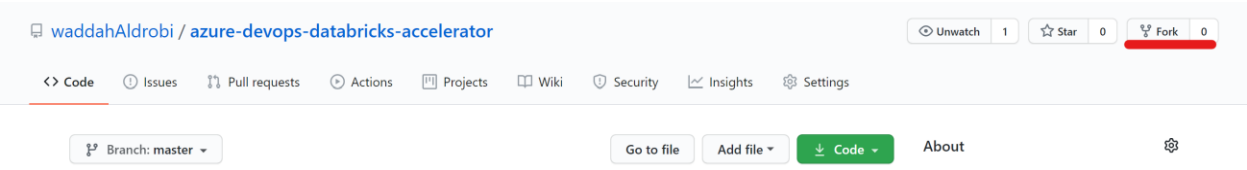

8. Creating an Azure DevOps Pipeline

1. Login into https://dev.azure.com/
2. Navigate to the project created earlier
3. In the project, navigate to *Pipelines*
4. Click *Create Pipeline*
5. Select *GitHub*
6. Select the GitHub repository from the previous step (follow any consequent instructions)
7. Click the arrow next to *Run* and click *Save*

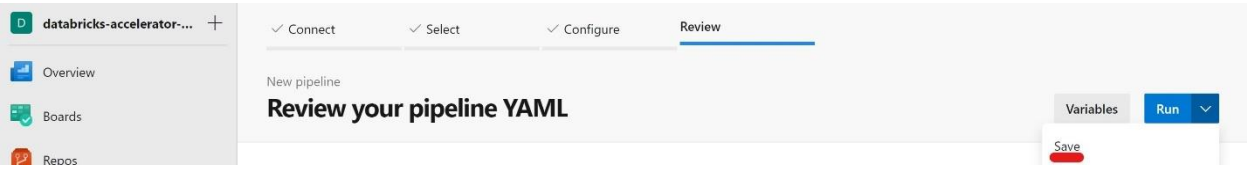

9. Retrieving Databricks Personal Access Token and Region URL

1. Sign in to the Azure portal
2. Navigate to Resource Groups and select the resource group created earlier
3. Select the Databricks resource (*accelerator-databricks)*
4. Click on *Launch Workspace*
5. In the top right corner, select the name of the workspace, then select *User Settings*

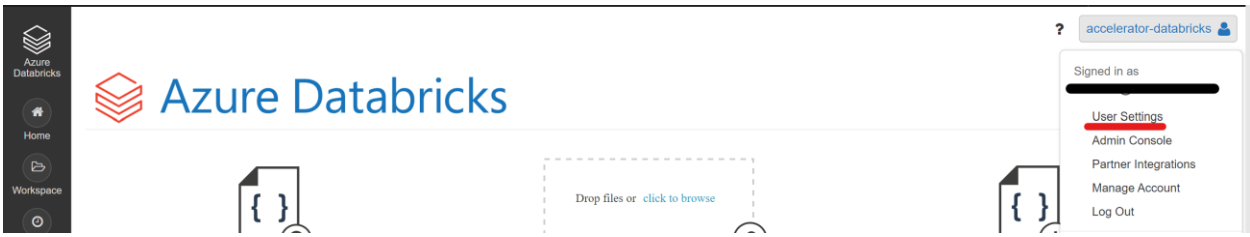

6. Under *Access Tokens,* click *Generate Token*
7. Fill in the required fields and click *Generate*
8. **Copy the token and store it securely.** (Should be in the form *dapi****) and click *Done*
9. Copy the part of URL in your browser that's in the form *https://***.azuredatabricks.net*

By now you should have the following hash values (necessary for the next step):

- The personal access token from Azure Databricks
- The workspace region URL

10. Adding the Personal Access Token to Key vault

1. Go to Azure Portal
2. Go to the Resource Group and open *acceleratorkeyvault* (or what you name the resource)
3. On the left menu, under Settings, click *Secrets*
4. Click *+ Generate/Import*
5. Fill in the fields.
    a. **Note:** the key must be **CSE-DEVELOP-PAT**, otherwise it will need to be changed in the pipeline's YAML file
    b. Set the *Value* as the personal access token has value retrieved earlier
6. Click *Create*

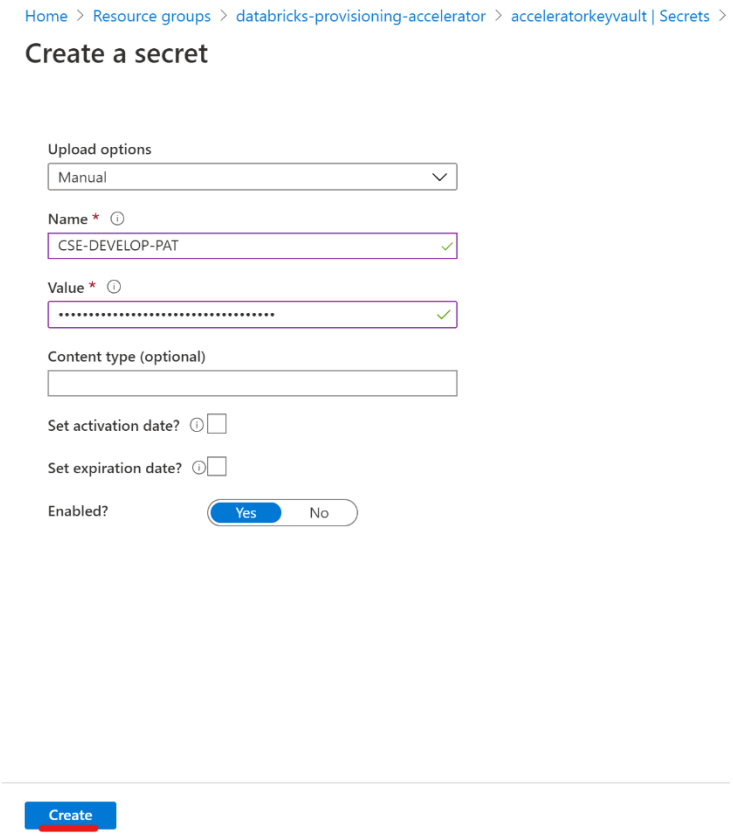

11. Adding Variables to the Pipeline

1. Login into https://dev.azure.com/
2. Navigate to the project created earlier
3. Navigate to *Pipelines* and select the pipeline created earlier
4. In the top right corner, click *Edit*
5. In the top right corner, click *Variables*

Note: the following variables MUST have the same names, otherwise they will need to be changed in the pipeline's YAML file

6. Click *New Variable*
    a. Set *Name* as **WORKSPACE-REGION-URL**
    b. Set *Value* as the workspace region URL copied earlier

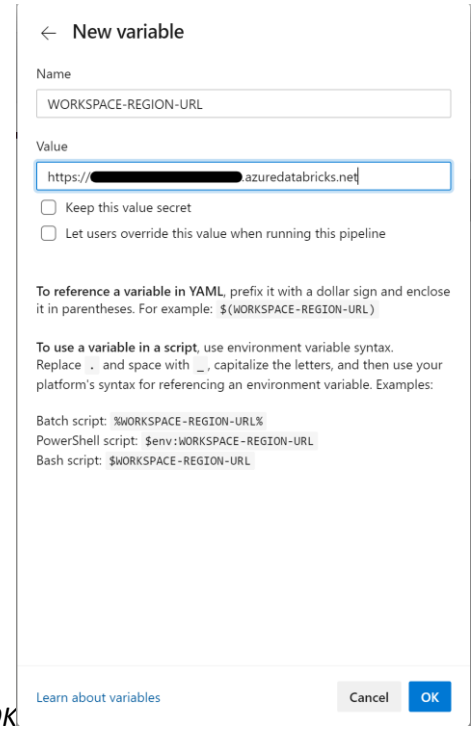

7. Click *OK*
8. Click on the *+* next to the search bar
    a. Set *Name* as **NOTEBOOK_NAME**
    b. Set *Value* as **mount.py**
9. Click OK
10. Click on the *+* next to the search bar
    a. Set *Name* as **SECERET-SCOPE-NAME**
    b. Set *Value* as **accelerator-secret-scope**
11. Click OK
12. Click on the *+* next to the search bar
    a. Set *Name* as **SCOPE-KEY-NAME**
    b. Set *Value* as **accelerator-secret**
13. Click on *OK, then* Click on *Save*

12. Adding a Variable Group containing Key Vault variables

1. Using the menu bar on the left
2. Under *Pipelines,* navigate to *Library*
3. Click *+ Variable group*
4. T must be named **keyvault_variable_group**, else edit the YAML file
5. Click on the toggle *Link secrets from an Azure key vault as variables*
6. Select the service connection created earlier (accelerator-service-connection)
7. Select the Key vault name created earlier (acceleratorkeyvault)
8. When an error is prompted, copy the PowerShell script provided
9. Navigate to [Azure portal]
10. Click on the cloud shell button as shown below

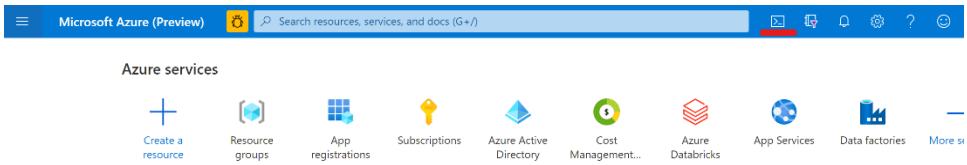

11. Make sure that a PowerShell environment is selected. Paste the script copied earlier and run it.
12. Follow the steps to authenticate.
13. To confirm, navigate to Resource Groups > <Project's resource group> > <Your Key Vault instance> > Access Policies. You should find your app's registration name under Applications.
14. Return to Azure DevOps, and click on the refresh button next to *Key vault name*
15. Under *Variables* click *+ Add* and select both *accelerator-storage-key* and *CSE-DEVELOP-PAT*
16. Click on *Save* (on the top bar)

13. Linking the Variable Group to the Pipeline

1. In Azure DevOps, click on *Pipelines*
2. Click on the pipeline created earlier
3. Click on *Edit* (Authorize if necessary)
4. Click on options (…)

5. Click on *Triggers*
6. Click on *Variables*
7. Click on *Variable Groups*
8. Click on *Link variable group*
9. Select the variable group created earlier
10. Click *Link*

14. Running the Pipeline

(Note: If using Azure Datalake Gen 2, the content of mount.py should be changed. For more info, refer to https://docs.microsoft.com/en-us/azure/databricks/data/data-sources/azure/azure-datalake-gen2#--mount-an-azure-data-lake-storage-gen2-account-using-a-service-principal-and-oauth-20)

1. In Azure DevOps, click on *Pipelines*
2. Click on the pipeline created earlier
3. Click on *Run pipeline*
4. Click on *Run*

For the first time running the pipeline, permission to the service connection might be required:

5. Click on *Job*
6. Click on *View*
7. Click on *Permit*
8. Click on *Permit* to the popup

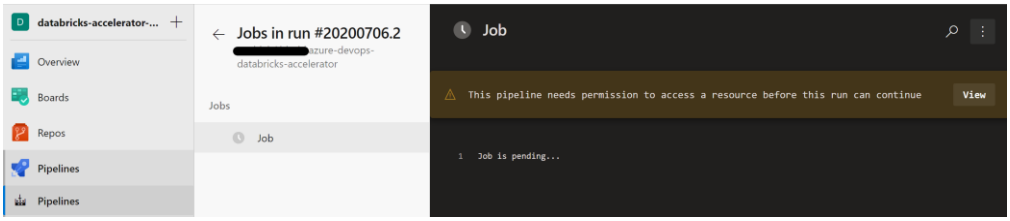

If everything has run successfully, the following is the result of pipeline's run:

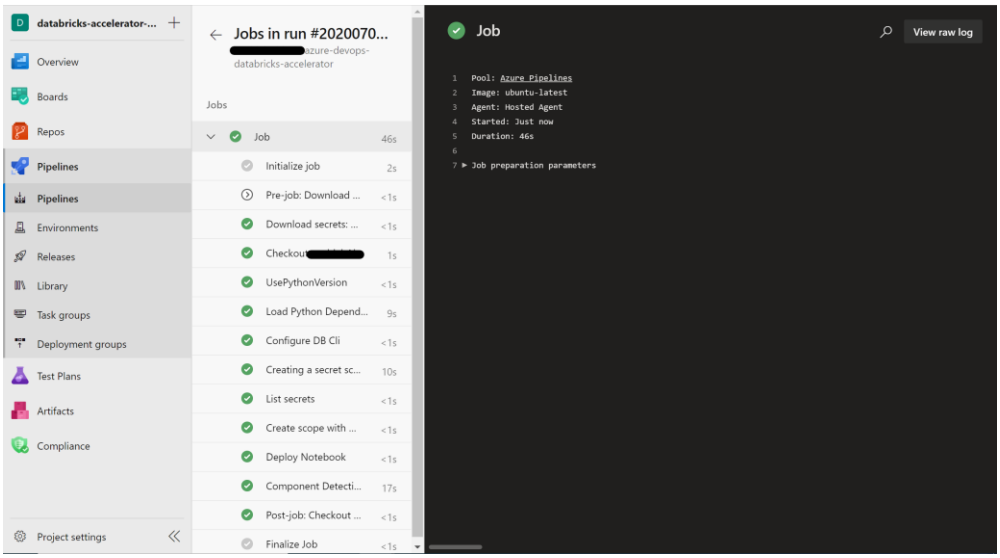

15. Mounting the Databricks Notebook to Azure Storage

1. Sign in to the [Azure portal]
2. Navigate to Resource Groups and select the resource group created earlier
3. Select the Databricks resource (*accelerator-databricks)*
4. Click on *Launch Workspace*
5. Click *Home* from the left menu bar

6. Under *Workspace,* click on *Shared*
7. Click on *mount-deployed,* you should see the following:

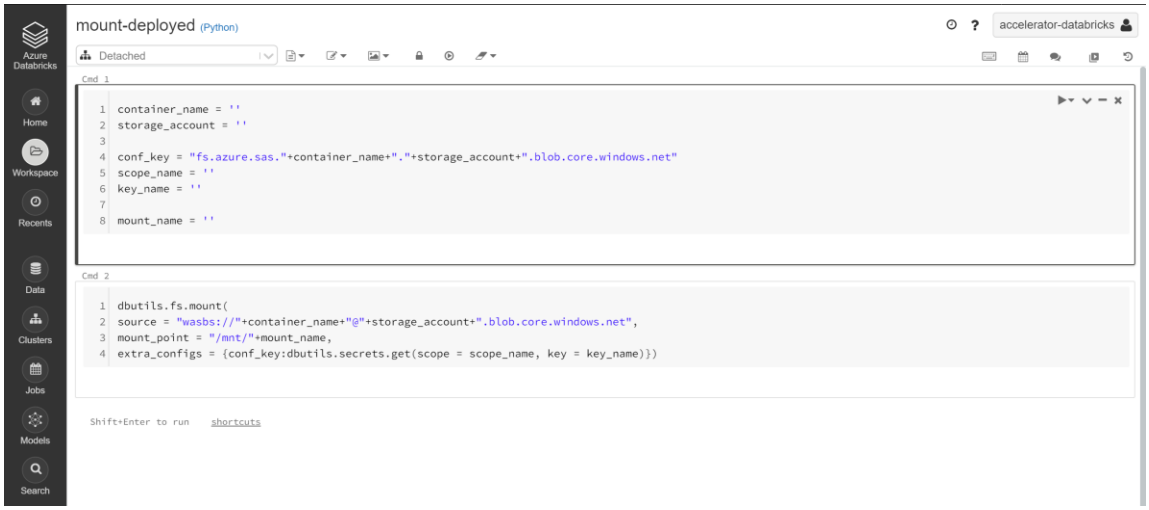

8. Click on *Detached* and click on *Create a Cluster* (or choose a cluster if available)
9. Fill in the require fields and click *Create Cluster*

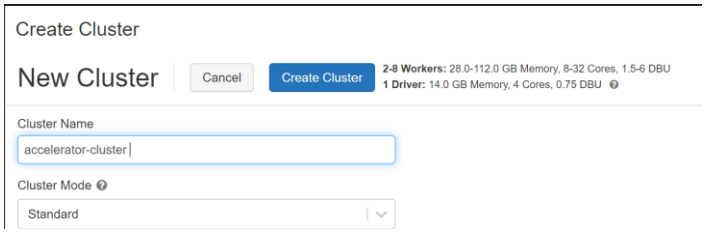

10. Wait for the cluster to run (take some time, about 5 minutes)
11. Return to the *mount-deployed* notebook
12. Click on *Detached* and attach the cluster created earlier
13. Fill in the variables with empty strings and run both cells (the output should be True)